

Chapter 5: Mutual Exclusion and Semaphores

CSCI 3753 Operating Systems

William Mortl, MS and Rick Han, PhD

University of Colorado at Boulder




Recap

- Inter-Process Communication (IPC)
 - Shared Memory
 - Message Passing
 - Unix-domain sockets, Internet sockets, pipes, signals
- Chapter 5: Synchronization
 - Race conditions occur when 2 processes/threads try to access the same shared variable
 - Example



Synchronization

Bounded Buffer



```
while(1) {  
    while(counter==MAX);  
    buffer[in] = nextdata;  
    in = (in+1) % MAX;  
    counter++;  
}
```

```
while(1) {  
    while(counter==0);  
    getdata = buffer[out];  
    out = (out+1) % MAX;  
    counter--;  
}
```

Producer writes new data into
buffer and increments counter

counter
updates
can conflict!

Consumer reads new data from
buffer and decrements counter



A Race Condition Example

- Let brackets [value] denote local value of counter in either the producer or consumer's process. counter=5 initially.

// counter++

(1) reg1 = counter; [5]

(3) reg1 = reg1 + 1; [6]

(5) counter = reg1; [6]

// counter--;

(2) reg2 = counter; [5]

(4) reg2 = reg2 - 1; [4]

(6) counter = reg2; [4]

- Counter should be 5 with 1 producer and 1 consumer, but counter = 4! Reversing steps (5) and (6) sets counter=6
- Undesirable and unpredictable *race condition*
- Basic Problem: unprotected access to a shared variable

(counter)



Critical Section

- Some kernel data structures could be subject to race conditions, e.g. access to list of open files
- Kernel developer must ensure that no such race conditions occur
- User or kernel developer identifies *critical sections* in code where each process accesses shared variables
 - access to critical sections is controlled by special *entry* and *exit* code

```
while(1) {
```

```
    entry section
```

```
        critical section (manipulate common var' s)
```

```
    exit section
```

```
        remainder section code
```

```
}
```



Critical Sections Need Mutual Exclusion

- How do we protect access to critical sections?
 - want to prevent another process from executing while current process is modifying this variable - *mutual exclusion*
- Approaches to mutual exclusion:
 - User space – see Peterson's solution Chapter 5.2
 - Hardware-based - disable interrupts
 - before entering critical section, disable interrupts
 - after exiting critical section, reenale interrupts
 - This provides mutual exclusion



Disabling Interrupts

```
shared int counter;
```

Code for p_1

```
disableInterrupts();  
counter++;  
enableInterrupts();
```

Code for p_2

```
disableInterrupts();  
counter--;  
enableInterrupts();
```

Achieves mutual exclusion:

- While P1 is in its critical section, P2 cannot enter, and vice versa
- Only when the first process is finished by enabling interrupts can another process attempt to enter the critical section
- modify the shared counter



Disabling Interrupts (2)

```
shared int counter;
```

Code for p_1

```
disableInterrupts();  
counter++;  
enableInterrupts();
```

Code for p_2

```
disableInterrupts();  
counter--;  
enableInterrupts();
```

Drawbacks?

- Interrupts could be disabled indefinitely
 - Due to a bug, e.g. user forgets to reenables interrupts
 - Unintentionally due to a loop



Disabling Interrupts (3)

```
shared int counter;
```

Code for p_1

```
disableInterrupts();  
counter++;  
enableInterrupts();
```

Code for p_2

```
disableInterrupts();  
counter--;  
enableInterrupts();
```

Drawbacks?

- Interrupts could be disabled a long time
 - Interrupts can be disabled too long, e.g. if the critical section has many lines of code
 - Can prevent useful I/O from being processed
 - Can prevent useful progress on other processes



Locks

- Goal: achieve mutual exclusion without the drawbacks of fully disabling interrupts
- Create a variable/flag called a *lock* to
 - The lock's value indicates whether the critical section (or access to its protected variables) is busy
 - Unset the lock or flag when the critical section is done
- Locks are often called *mutexes*



Lock Semantics

- Each task calls `Acquire(lock)` before entering its critical section.
 - If a task successfully grabs the lock, this task can proceed to execute its critical section
 - If the lock is not available, the calling task blocks. So only 1 task at a time has the lock.
 - Multiple tasks can block waiting on a lock, forming a FIFO queue
- A task calls `Release(lock)` when done, so other processes can enter their critical sections too
 - 1 waiting task is woken (based on order) at each lock release



A Lock Example

```
shared int counter;  
shared Lock lock;
```

Code for p_1

```
Acquire(lock) ;  
...  
counter++;  
...  
Release(lock) ;
```

Code for p_2

```
Acquire(lock) ;  
...  
counter--;  
...  
Release(lock) ;
```

Advantage: Achieves mutual exclusion without blocking all processes

- Other tasks are blocked from executing their critical sections
- But if the owner of the lock is context-switched out, other tasks can keep executing (provided they have no interest in the lock)



A Lock Example

```
shared int counter;  
shared Lock lock;
```

Code for p_1

```
Acquire(lock) ;  
...  
counter++;  
...  
Release(lock) ;
```

Code for p_2

```
Acquire(lock) ;  
...  
counter--;  
...  
Release(lock) ;
```

- Lock-based solutions only work if all processes participate and surround critical sections with entry and exit code to synchronize access to shared data
- Acquire and Release implemented as system calls



Locks

- Additional advantages:
 - user doesn't have control over disabling/reenabling interrupts – the OS does this.
 - user doesn't have to remember to reenale interrupts
 - So tasks with I/O that need to process interrupts can make progress
- Disadvantage:
 - user still must remember to release(lock), otherwise other tasks blocked on the lock will wait forever



A Flawed Lock Implementation

```
shared boolean lock = FALSE;  
shared int counter;
```

Code for p₁

```
/* Acquire the lock */  
while(lock){ no_op;}  
lock = TRUE;
```

```
/* Execute critical  
   section */  
counter++;
```

```
/* Release lock */  
lock = FALSE;
```

Code for p₂

```
/* Acquire the lock */  
while(lock){ no_op;}  
lock = TRUE;
```

```
/* Execute critical  
   section */  
counter--;
```

```
/* Release lock */  
lock = FALSE;
```

Testing of the lock using a while() is subject to a race condition

A Flawed Lock Implementation

```
shared boolean lock = FALSE;  
shared int counter;
```

Code for p₁

```
/* Acquire the lock */  
while(lock){ no_op;}  
lock = TRUE;
```

```
/* Execute critical  
   section */  
counter++;
```

```
/* Release lock */  
lock = FALSE;
```

- If P1 has just completed the *test* “while(lock)” and is just before the *setting* of “lock=TRUE”, P1 can get context switched out
- Then P2 executes and sees lock=FALSE in its while(lock) test, and also drops through to the lock=TRUE statement
- Now both P1 and P2 are executing in critical section - BAD

A Flawed Lock Implementation

```
shared boolean lock = FALSE;  
shared int counter;
```

Code for p₁

```
/* Acquire the lock */  
while(lock){ no_op;}  
lock = TRUE;
```

```
/* Execute critical  
   section */  
counter++;
```

```
/* Release lock */  
lock = FALSE;
```

- The problem: between *testing* the lock and *setting* the lock to FALSE, a task can be context-switched out
- Instead, want an uninterruptible or *atomic* test-and-set operation
 - Achieve by disabling interrupts
 - Achieved by hardware test-and-set instruction TS

A Correct Lock Implementation

```
Acquire(lock) {  
    disableInterrupts();  
  
    /* Wait for lock */  
    while(lock) {  
        enableInterrupts();  
        disableInterrupts();  
    }  
    lock = TRUE;  
    enableInterrupts();  
}
```

```
Release(lock) {  
    disableInterrupts();  
    lock = FALSE;  
    enableInterrupts();  
}
```

- Advantage: no race condition as test-and-set is atomic
- Disadvantage:
 - busy waiting to acquire lock (A busy waiting type of lock is also called a spinlock)
 - task is stuck in Acquire() and can't do other work



Test-and-Set Instruction

- CPU hardware provides an uninterruptible hardware instruction TS,
 - which atomically performs *both* the test and the set operations
 - called by the TestandSet() system call
- the boolean TestandSet() operation in the next slide is essentially a swap of values
 - The x86 CPU instruction set contains atomic instructions such as XCHG that are essentially swap statements
 - Can use atomic XCHG to implement spinlocks



A Test-and-Set Lock Implementation

```
shared boolean lock = FALSE;  
shared int counter;
```

Code for p_1

```
/* Acquire the lock */  
while (TestandSet(&lock)) ;  
  
/* Execute crit sect */  
counter++;  
  
/* Release lock */  
lock = FALSE;
```

Code for p_2

```
/* Acquire the lock */  
while (TestandSet(&lock)) ;  
  
/* Execute crit sect */  
counter--;  
  
/* Release lock */  
lock = FALSE;
```

```
boolean TestandSet(boolean *target) {  
    boolean rv = *target;  
    *target = TRUE;  
    return rv;  
}
```

This entire instruction
(sequence) is atomic. The
hardware enforces atomicity.



A Test-and-Set Lock Implementation

- Mutual exclusion is achieved - no race conditions
 - If one process X tries to obtain the lock while another process Y already has it, X will wait in the loop
 - If a process is testing and/or setting the lock, no other process can interrupt it
- Efficient:
 - The system is exclusively occupied for only a short time - the time to test and set the lock, and not for entire critical section
 - typically only about 10 instructions



A Test-and-Set Lock Implementation

- Advantage:
 - Compared to spinning inside the Acquire() system call, this approach spins in user space
 - Allows the user task to do other execution while polling the TestandSet()
- Disadvantage:
 - requires user to remember to put in while()
 - while() means busy-waiting, i.e. this is a form of spinlock too except in user space
 - requires special hardware support in the form of a low level atomic machine instruction TS



Blocking on Locks

- Simplicity of spinlocks:
 - don't require any other OS constructs, such as a queue of tasks blocked on the lock
 - If multiple tasks are blocked on the lock, they all spin
 - Can be useful if blocking time is short
- Locks can be augmented with queues to hold blocked tasks
 - OS scheduler must be involved
 - Useful when blocking time is long
 - We'll see an example of this with semaphores



Semaphores

- more general solution to mutual exclusion proposed by Dijkstra
- Semaphore S is an integer variable that, apart from initialization, is accessed only through 2 standard atomic operations
 - P(), also called wait()
 - somewhat equivalent to a test-and-set, but also *decrements* the value of S
 - V(), also called signal()
 - *increments* the value of S
 - OS provides ways to create and manipulate semaphores atomically



Semaphores

- Pseudo-code for classic semaphore
 - its value can't go below zero, i.e. classic semaphore is non-negative

“atomic”	{	<code>P(S) {</code>		<code>V(S) {</code>	}	atomic
		<code> while(S<=0)</code>		<code> S++;</code>		
		<code> {wait or no op;}</code>		<code>}</code>		
		<code> S--;</code>				
	}					

P() is “atomic” in the sense that it tests S atomically, and if $S \leq 0$, then the process calling P() will *relinquish control*. Otherwise, the process continues forward and decrements S atomically.



A Semaphore Implementation

- Based only disabling and reenabling of interrupts
 - semaphores can also be implemented using TestandSet() instructions

```
P(S) {  
    disableInterrupts();  
    while(S<=0) {  
        enableInt();  
        disableInt();  
    }  
    S--;  
    enableInt()  
}
```

```
V(S) {  
    disableInt();  
    S++;  
    enableInt()  
}
```

Also a form of spinlock
S not decremented unless it is positive

disable interrupts
because S++ might
cause race condition



Mutual Exclusion with Semaphores

```
Semaphore S = 1; // initial value of semaphore is 1
int counter;      // assume counter is set correctly somewhere in
                  // code
```

Process P1:

```
P(S);
    // execute crit sect
    counter++;
V(S);
```

Process P2:

```
P(S);
    // execute crit sect
    counter--;
V(S);
```

- Both processes atomically P() and V() the semaphore S, which enables mutual exclusion on critical section code, in this case protecting access to the shared variable counter



Mutual Exclusion with Semaphores (2)

```
Semaphore S = 1; // initial value of semaphore is 1
int counter;      // assume counter is set correctly somewhere in
                  // code
```

Process P1:

```
P(S);
// execute crit sect
counter++;
V(S);
```

Process P2:

```
P(S);
// execute crit sect
counter--;
V(S);
```

- The 1st process that calls P() on semaphore S will check if S==0 (it is not), and then will atomically decrement S to 0
- The next process calling P(S) blocks on semaphore S since S=0, hence mutual exclusion is achieved
- When 1st process is done, it calls V(S), atomically incrementing S to 1



Mutual Exclusion with Semaphores (3)

```
Semaphore S = 1; // initial value of semaphore is 1
int counter;      // assume counter is set correctly somewhere in
                  code
```

Process P1:

```
P(S);
    // execute crit sect
    counter++;
V(S);
```

Process P2:

```
P(S);
    // execute crit sect
    counter--;
V(S);
```

- After a V(S), all waiting processes are busy-waiting, so there is a race to see which one calls P(S) first



Binary Semaphores

- The previous example showed how to use the semaphore as a *binary* semaphore
 - its initial value was set to 1
 - a binary semaphore is also called a *mutex lock*, i.e. it can be used to provide mutual exclusion on some piece of critical code
 - Let's define a binary semaphore as a semaphore whose value can't exceed 1, even if many processes keep `signal()`'ing the semaphore
 - Additional logic would have to be added to the counting semaphore in its `V()/signal()` function to cap its maximum value to 1



Counting Semaphores

- A semaphore can also be used more generally as a *counting* semaphore
 - its initial value is n , e.g. $n=10$
 - the value of the semaphore could be used to keep track of the number of instances of a finite resource that are still available
 - We'll see an example later using counting semaphores to solve the general producer/consumer bounded buffer problem



Enforcing Order with Semaphores

- Enforce *order* of access between 2 processes P1 and P2
 - P1 contains code C1 and P2 contains code C2
 - **Want to ensure that code C1 executes before code C2**
 - Use semaphores to synchronize the order of execution of the two processes

Semaphore S=0; // initial value of semaphore = 0

Process P1:

```
C1;           // execute C1
signal(S);    // V() semaphore
```

Process P2:

```
wait(S);     // P() semaphore
C2;          // execute C2
```



Enforcing Order with Semaphores (2)

- if P1 executes 1st, then
 - C1 will execute 1st, then P1 V()'s semaphore, adding 1 to its value
 - Later, when P2 executes, it will call wait(S), which will decrement the semaphore to 0 – no waiting - followed by execution of C2
 - Thus C1 executes before C2

Semaphore S=0; // initial value of semaphore = 0

Process P1:

```
C1;           // execute C1
signal(S);    // V() semaphore
```

Process P2:

```
wait(S);     // P() semaphore
C2;          // execute C2
```



Enforcing Order with Semaphores (3)

- If P2 executes 1st, then
 - P2 blocks on semaphore (=0), so C2 will not be executed yet
 - Later, when P1 executes, it runs C1, then V()'s the semaphore
 - This awakens P2, which then executes C2
 - Thus C1 executes before C2

Semaphore S=0; // initial value of semaphore = 0

Process P1:

```
C1;           // execute C1
signal(S);    // V() semaphore
```

Process P2:

```
wait(S);     // P() semaphore
C2;          // execute C2
```



A Revised Semaphore Definition

```
typedef struct {  
    int value;  
    struct process *list;  
} semaphore;
```

atomic {

```
P(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process  
        to S->list;  
        block();  
    }  
}
```

atomic {

```
V(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P  
        from S->list;  
        wakeup(P);  
    }  
}
```

- Efficiently sleep the process until it needs to be woken up by a V()/signal(), rather than spinlock



A Revised Semaphore Definition

- New definition allows a semaphore's value to be negative, because the decrement occurs before the test in $P()$
 - The absolute value of the semaphore's negative amount can be used to indicate the # of processes blocked on the semaphore
- Processes now yield the CPU if the semaphore's value is negative, rather than busy wait
 - If more than one process is blocked on a semaphore, then use a FIFO queue to select the next process to wake up when a semaphore is V'ed
 - Why is LIFO to be avoided?



Mutual Exclusion with Revised Semaphore

```
Semaphore S = 1; // initial value of semaphore is 1
int counter;      // assume counter is set correctly somewhere in
                  // code
```

Process P1:

```
P(S);
counter++;
V(S);
```

Process P2:

```
P(S);
counter--;
V(S);
```

Process P3:

```
P(S);
counter*=2;
V(S);
```

- The 1st process (say P2) that calls P() on semaphore S will set $S=0$ and move into its critical section
- The next process (say P3) calling P(S) decrements S to -1 and blocks on semaphore, hence mutual exclusion is achieved
- The next process (say P1) calling P(S) decrements S to -2 and blocks on S – *Note how $|S| = \# \text{ blocked processes!}$*



Mutual Exclusion with Revised Semaphore

```
Semaphore S = 1; // initial value of semaphore is 1
int counter;      // assume counter is set correctly somewhere in
                  // code
```

Process P1:

```
P(S);
counter++;
V(S);
```

Process P2:

```
P(S);
counter--;
V(S);
```

Process P3:

```
P(S);
counter*=2;
V(S);
```

- P2 finishes, V()'s the semaphore S, increasing its value to -1, signaling a process to unblock (P3).
- P3 finishes, V()'s the semaphore S, increasing its value to 0, causing a process to unblock (P1)
- P1 finishes, V()'s the semaphore S, increasing its value to original value of 1. No more processes to unblock.



Enforcing Order with Revised Semaphore

- If P1 hits its signal(S) first, before P2 hits wait(S), then
 - P1 will have executed C1 already, will then increment S to 1, and no process will be unblocked.
 - Later, P2 calls wait(S), decrements S from 1 to 0, and executes C2. Order is preserved: C1 executes before C2

Semaphore S=0; // initial value of semaphore = 0

Process P1:

```
C1;           // execute C1
signal(S); // V() semaphore
```

Process P2:

```
wait(S);    // P() semaphore
C2;          // execute C2
```



Enforcing Order with Revised Semaphore

- If P2 hits wait(S) before P1 hits signal(S), then
 - P2 will decrement S to -1, and block P2 on the semaphore.
 - Next, P1 calls signal(S) having executed code C1, incrementing S from -1 to 0 and unblocking P2.
 - P2 now executes C2. Order is preserved: C1 executes before C2

Semaphore S=0; // initial value of semaphore = 0

Process P1:

```
C1;           // execute C1
signal(S); // V() semaphore
```

Process P2:

```
wait(S);    // P() semaphore
C2;          // execute C2
```



Deadlock

- Semaphores provide synchronization, but can introduce more complicated higher level problems like *deadlock*
 - two processes deadlock when each wants a resource that has been locked by the other process
 - e.g. P1 wants resource R2 locked by process P2 with semaphore S2, while P2 wants resource R1 locked by process P1 with semaphore S1



Deadlock Example

```
Semaphore Q = 1;    // binary semaphore as a mutex lock
Semaphore S = 1;    // binary semaphore as a mutex lock
variable R1, R2;
```

Process P1:

Process P2:

P(S);	(step 1)	(step 2) P(Q);
P(Q);	(step 3)	(step 4) P(S);
modify R1 and R2;	<i>Deadlock!</i>	modify R1 and R2;
V(S);		V(Q);
V(Q);		V(S);

If steps (1) through (4) are executed in that order, then P1 and P2 will be deadlocked after statement (4) - verify this for yourself by stepping thru the semaphore values



Deadlock

- In the previous example,
 - Each process will sleep on the other process's semaphore
 - the V() signalling statements will never get executed, so there is no way to wake up the two processes from within those two processes
 - there is no rule prohibiting an application programmer from P()'ing Q before S, or vice versa - the application programmer won't have enough information to decide on the proper order
 - in general, with N processes sharing N semaphores, the potential for deadlock grows



Other Deadlock Examples

- A programmer mistakenly follows a P() with a second P() instead of a V(), e.g.

P(mutex)
critical section

P(mutex)

- This causes a self-deadlock!

- A programmer forgets and omits V(mutex). Can cause deadlock.

P1:
P(mutex)
critical section 1

P2:
P(mutex)
critical section 2

V(mutex)

- P2 calls P(mutex) and executes crit sect 2. Then P1 blocks on P(mutex), then P2 blocks on P(mutex)



Other Synchronization Errors

- A programmer forgets and omits $P(\text{mutex})$. Can violate mutual exclusion if $P(\text{mutex})$ is omitted.

P1:

$P(\text{mutex})$

critical section 1

$V(\text{mutex})$

P2:

critical section 2

$V(\text{mutex})$

- A programmer reverses the order of $P()$ and $V()$, e.g.

$V(\text{mutex})$

critical section <---- this violates mutual exclusion,

$P(\text{mutex})$

but is not deadlock

