

Chapter 8: Paging and Segmentation

CSCI 3753 Operating Systems
Prof. Rick Han



Recap: Paging


- OS maintains a page table for each process
- Given a logical address, MMU finds its logical page, then looks up physical frame in page table

Logical Address
Space

page 0
page 1
page 2
page 3
page 4

Page Table

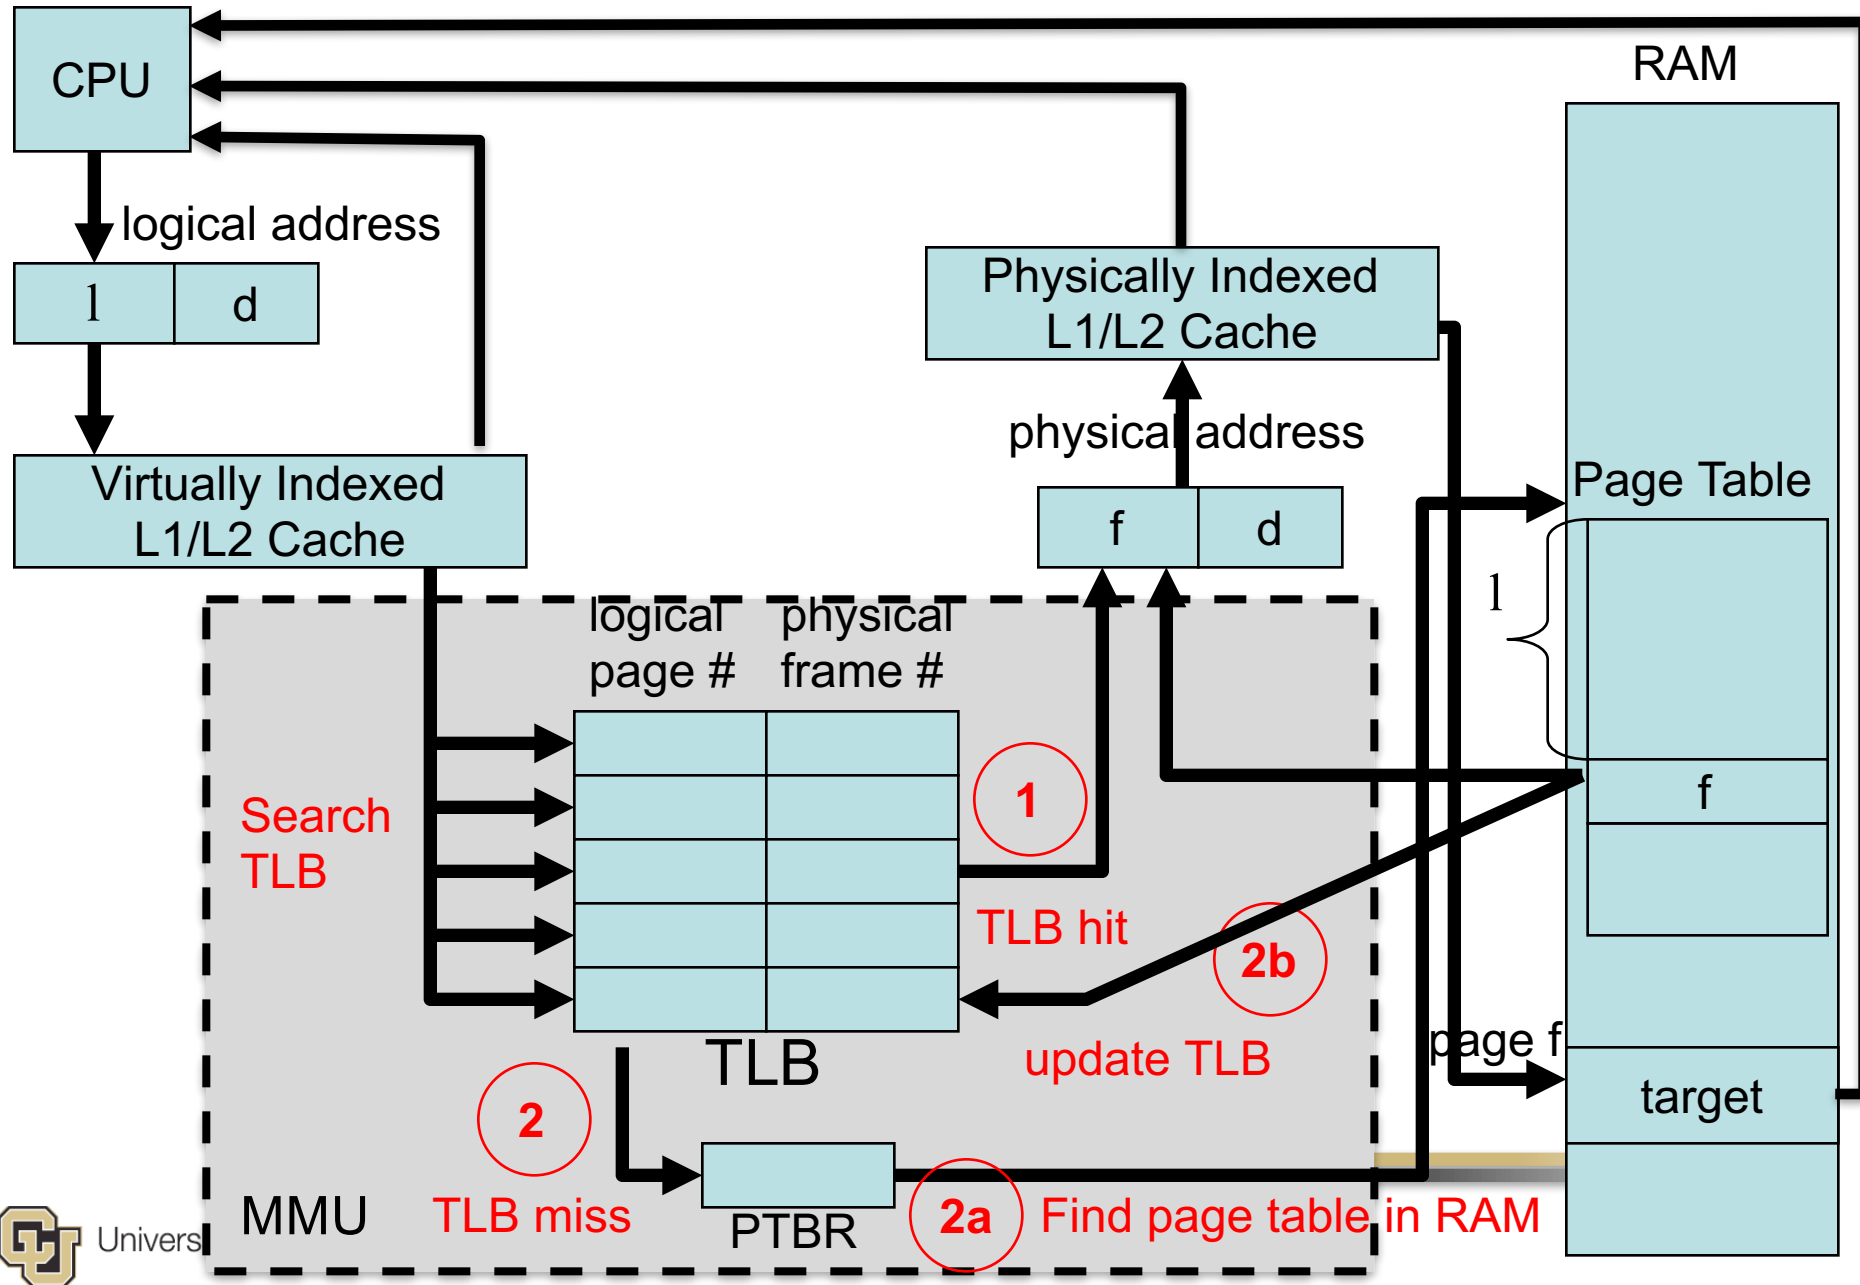
Logical page	Physical frame
0	1
1	4
2	3
3	7
4	0



could be allocated to
a new process P2

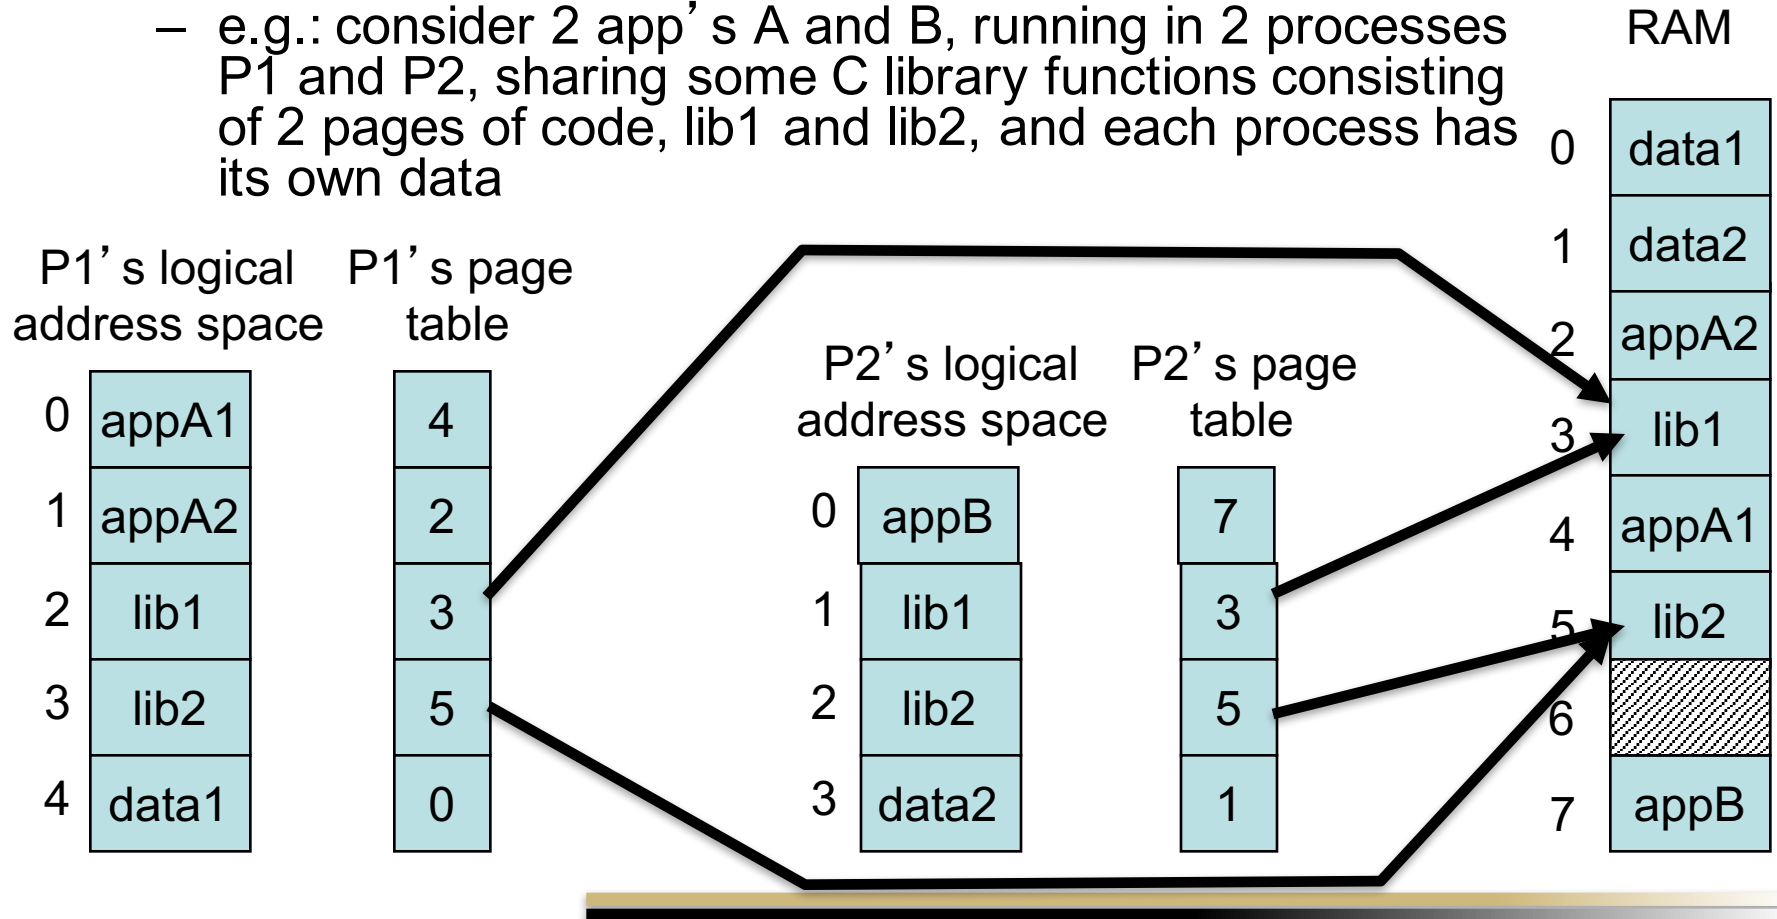


Recap: Paging with TLB and L1/L2 Caching



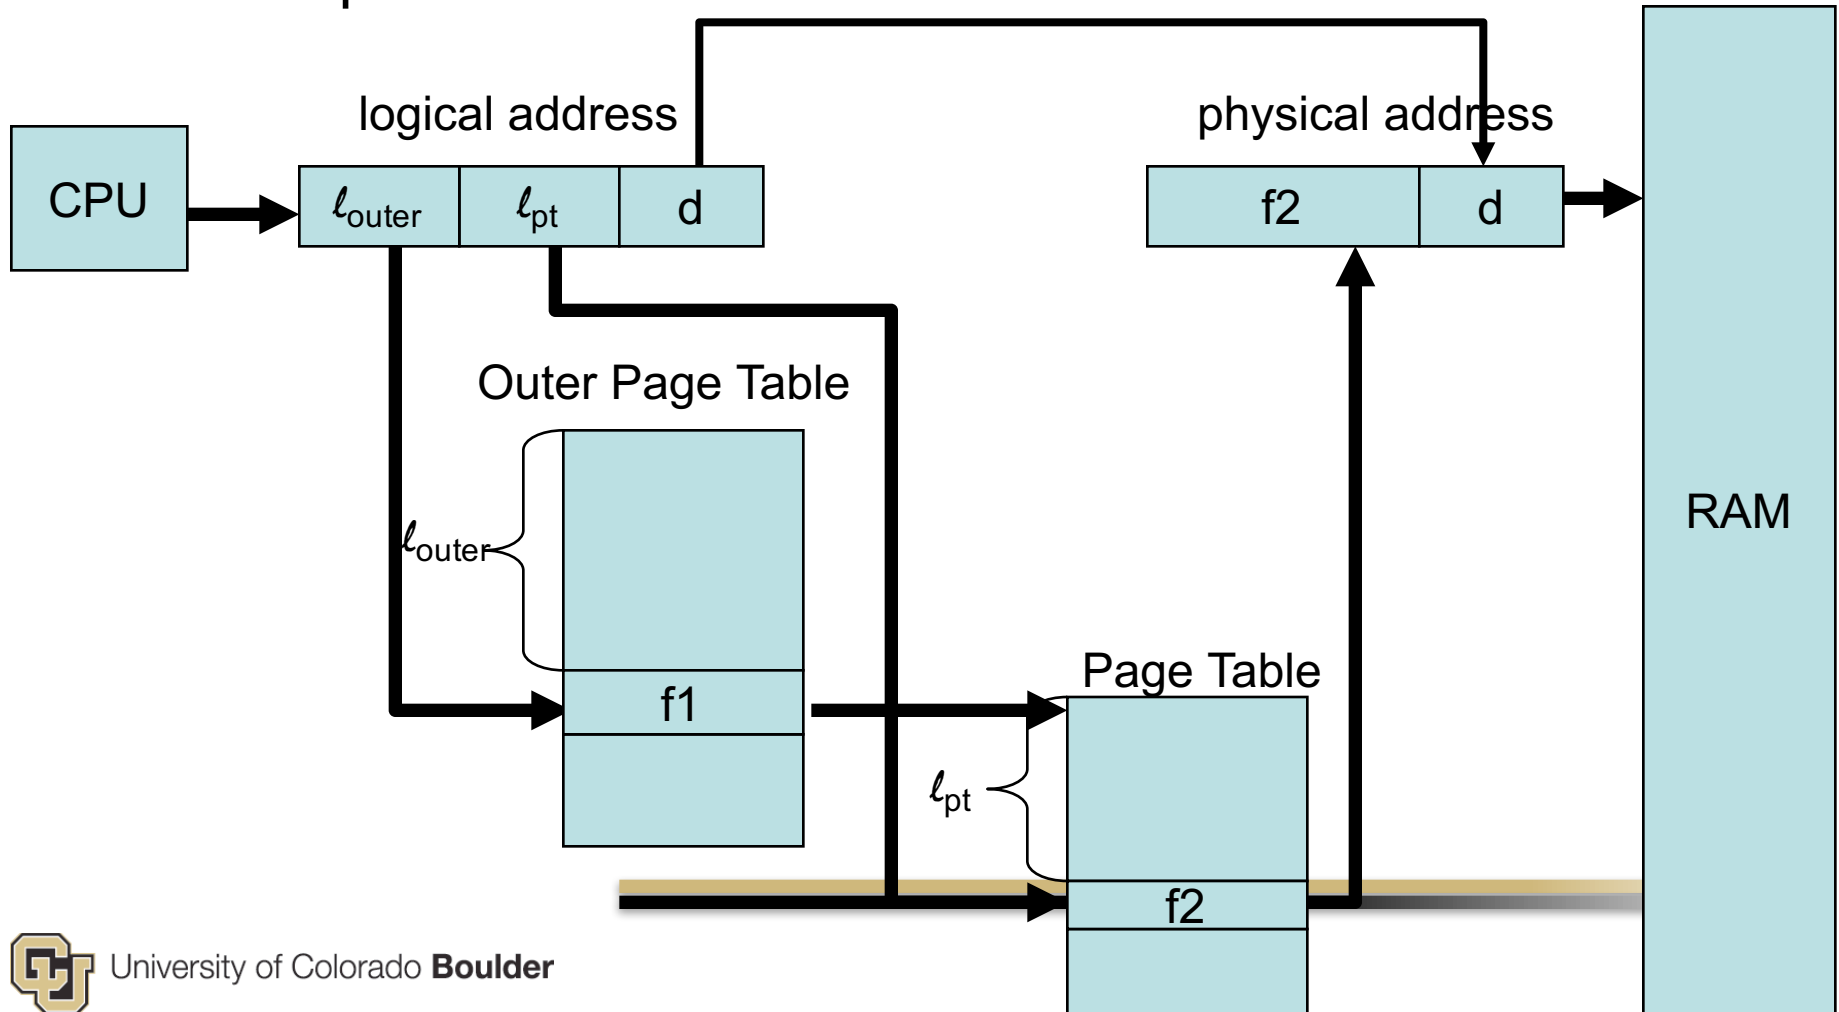
Recap: Shared Pages

- To achieve sharing of code and data, page tables can point to the *same* memory frames
 - e.g.: consider 2 app's A and B, running in 2 processes P1 and P2, sharing some C library functions consisting of 2 pages of code, lib1 and lib2, and each process has its own data



Recap: Hierarchical Paging

- Hierarchical (2-level) paging divides the logical address into 3 parts:



Sparse Page Tables

- Size of page tables may be large – millions of entries
- Few entries in each page table may be filled in
 - e.g. stack may never grow large enough to occupy a large fraction of virtual memory
 - e.g. heap may never grow large enough to occupy a large fraction of virtual memory
- Hence, there can be much “wasted” space in page tables
 - Still have to allocate space in memory for these unused page of the page table



Inverted Page Table

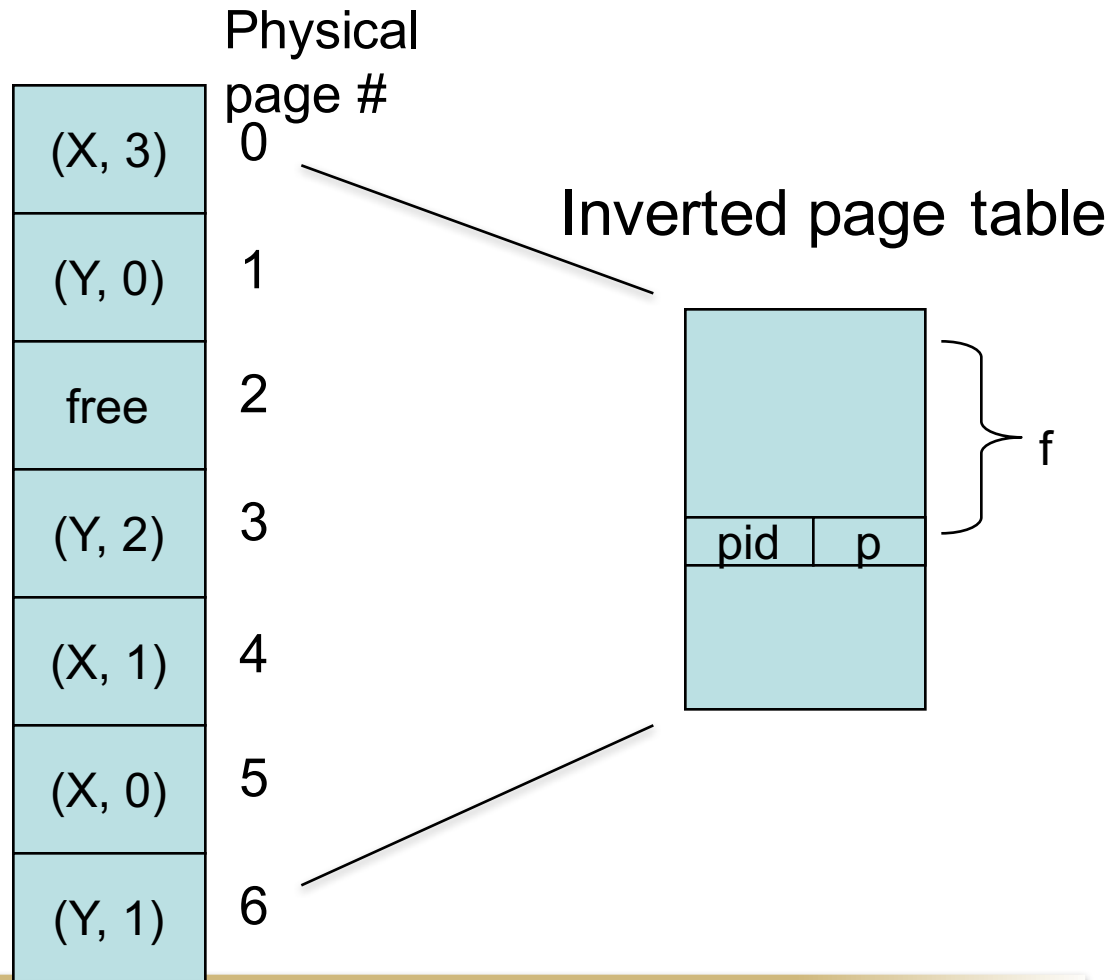
- Solution: use inverted page table (IPT)
 - Have only one page table for all of memory, rather than one for each process
 - this saves on memory
 - Each entry in the inverted page table lists which process owns a physical frame f , and what logical page # p is stored in that physical frame
 - each entry in the IPT lists a process id pid and a logical page p , namely $IPT[f] = \langle pid, p \rangle$
 - thus the index into an inverted page table is the physical frame # f
 - compare to a page table, whose index is the logical page # p



Inverted Page Table

Main memory

- Each physical page stores a logical page owned by some process
- Process X's logical page 3 is stored in physical page 0:
 - e.g. (X,3) in 0
- Collect all such assignments into a separate inverted page table array

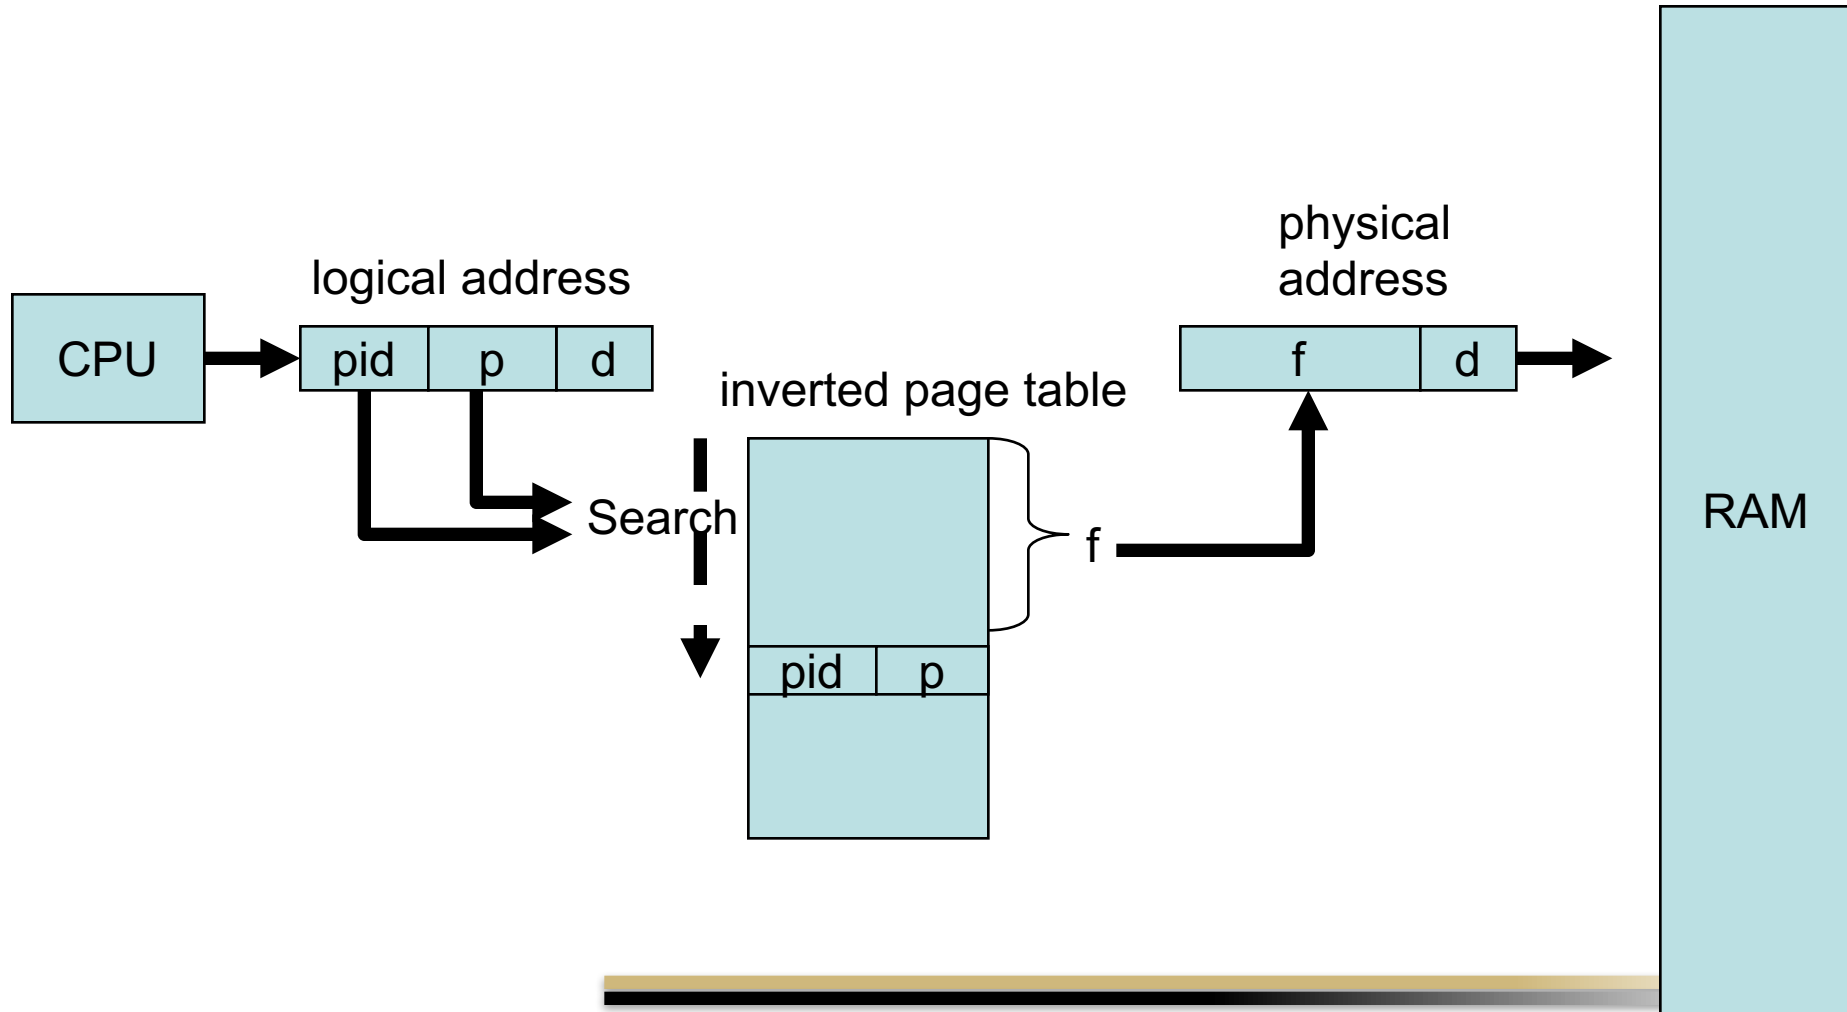


Inverted Page Table

- Still want to use an inverted page table to map a logical page to a physical frame
 - the IPT stores an array of $\langle \text{pid}, p \rangle$ pairs,
 - given a pid and logical page p , *search* through IPT to find an entry that matches
 - this can be slow - a hash table can speed this up (later slides)
- another drawback: it's hard to implement shared memory pages with IPTs
 - since only one process owns each physical frame,
 - whereas in shared memory multiple processes can use the same code in a physical frame



Inverted Page Table



Hashed Page Table

- A hash table allows you to search fewer entries in either an inverted page table, or a regular page table
- A hash function h is a many-to-one function
 - example: take the most significant byte of a 32-bit address, i.e. $h(32\text{-bit addr}) = \text{MSB}$
 - there may be more than one input value that maps onto the same output value, because it is many-to-one
 - if h takes the MS bit, then $h(10) = 1 = h(11)$, so both 10 and 11 map to the same hashed value



Hashed Page Table


- for an IPT:
 - input values $\langle \text{pid}, p \rangle$ into h and return a hashed value $h(\text{pid}, p) = H$.
 - Then index into the hash table with value H .
 - Each entry of the hash table contains a linked list, in case there are other input values that map to the same hashed value H .
 - Search through this linked list.
 - This is faster than searching the whole IPT linearly from the beginning



Hashed Inverted Page Table

Inverted
page table

Linked lists



- Hash takes as input $\langle \text{pid}, p \rangle$
- Note that empty entries of the inverted page table never need to be hashed




Hashed Page Table

- For a regular page table,
 - Hashing is useful for efficiently tracking assigned parts of very large *sparse* logical address spaces
 - example: typical heap and stack may be quite sparse with many empty logical addresses
 - Apply hashing with linked lists to track all logical addresses
 - Only assigned addresses will invoke the hash and will be tracked as indexes in a hash table
 - unassigned pages never invoke hash, so no entries to track!
 - hash table, even with a linked list per entry to search, is seen as smaller and faster



Hashed Page Table



- Hash takes as input the logical page # p
- Note that empty entries of the page table never need to be hashed



Segmentation

- An alternative solution for external fragmentation is to employ variably sized *segments*
 - Contrast to fixed-sized pages
 - Divide a process into variably sized segments that are organized according to some logical criteria
 - a process has code, data, stack, and heap - each of these could be a separate segment
 - a process could subdivide its code into functional segments
 - e.g. a Web server could divide its functional components into a networking segment, a database interface segment, and a dynamic page composition segment, etc.

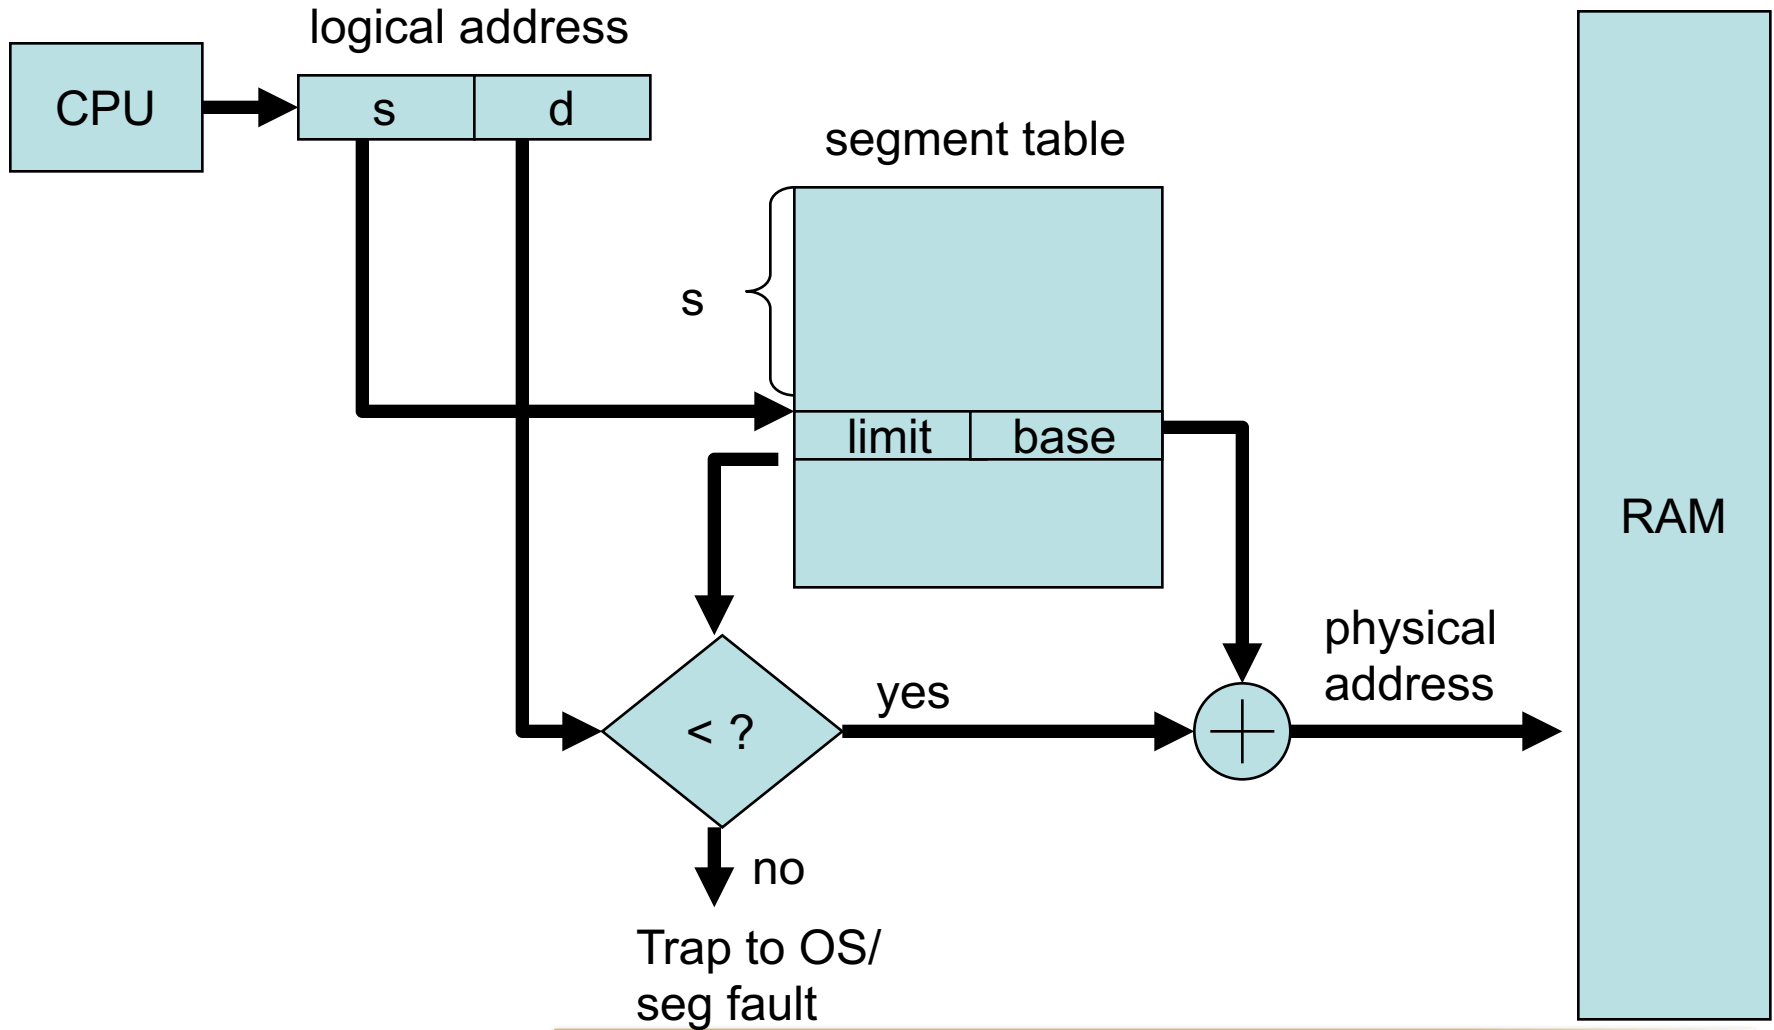


Segmentation

- Each instruction in a segment has a logical address = <segment #, offset>
- Need a segment table to keep track of where each segment is mapped in main memory
- Segments are variably sized
 - Each entry of the segment table needs a base and limit field to place in the base and limit registers



Segmentation



Segmentation

- x86 Pentium supports pure segmentation and segmentation with paging
 - Segmentation with paging means that each segment is divided into pages
- Linux on x86 Pentium uses segmentation sparingly
- Drawback of segmentation:
 - fragmentation of free space in RAM becomes even more complex and hard to manage

